

Practical Constructions and New Proof Methods for Large Universe Attribute-Based Encryption

Yannis Rouselakis
The University of Texas at Austin
jrous@cs.utexas.edu

Brent Waters^{*}
The University of Texas at Austin
bwaters@cs.utexas.edu

ABSTRACT

We propose two large universe Attribute-Based Encryption constructions. In a large universe ABE system any string can be used as an attribute and attributes need not be enumerated at system setup. Our first construction establishes a novel large universe Ciphertext-Policy ABE scheme on prime order bilinear groups, while the second achieves a significant efficiency improvement over the large universe Key-Policy ABE system of Lewko-Waters and Lewko. Both schemes are selectively secure in the standard model under two “ q -type” assumptions similar to ones used in prior works. Our work brings back “program and cancel” techniques to this problem and aims in providing practical large universe ABE implementations.

To showcase the efficiency improvements over prior constructions, we provide implementations and benchmarks of our schemes in Charm; a programming environment for rapid prototyping of cryptographic primitives. We compare them to implementations of the only three published constructions that offer unbounded ABE in the standard model.

Categories and Subject Descriptors

E.3 [Data]: Data Encryption—*Public key cryptosystems*

General Terms

Security; Theory

^{*}Supported by NSF CNS-0915361 and CNS-0952692, CNS-1228599 DARPA through the U.S. Office of Naval Research under Contract N00014-11-1-0382, DARPA N11AP20006, Google Faculty Research award, the Alfred P. Sloan Fellowship, Microsoft Faculty Fellowship, and Packard Foundation Fellowship. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the Department of Defense or the U.S. Government.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

CCS'13, November 4–8, 2013, Berlin, Germany.

Copyright is held by the owner/author(s). Publication rights licensed to ACM.

ACM 978-1-4503-2477-9/13/11 ...\$15.00.

<http://dx.doi.org/10.1145/2508859.2516672>.

Keywords

attribute-based encryption; large universe; unbounded; q -type assumptions; ciphertext-policy; key-policy; Charm

1. INTRODUCTION

Traditionally, public key encryption schemes provided any user with the ability to share data with another specific user in a private manner. However, in many applications we would like to have the additional capability to encrypt data for a set of users according to a specific policy on their credentials. For example, one might want to store data in a public server such that only parties with credentials of specific forms are able to decrypt. Instead of encrypting the data once for each party it would be beneficial to be able to encrypt only once for all desired parties. This encryption notion, called *Attribute-Based Encryption* (ABE), was introduced by Sahai and Waters [36]. In this setting, each user possesses a set of attributes/credentials and a secret key that corresponds to these credentials. The encrypting party can define any Boolean formula on the possible attributes and a user can decrypt if and only if his attribute set satisfies the Boolean formula.

Several attribute-based constructions have been presented since then (see related work below). A common classification property is whether a system is a “small universe” or “large universe” constructions. In “small universe” constructions the size of the attribute space is polynomially bounded in the security parameter and the attributes were fixed at setup. Moreover, the size of the public parameters grew linearly with the number of attributes. In “large universe” constructions, on the other hand, the size of the attribute universe can be exponentially large, which is a desirable feature.

Achieving the large universe property can be challenging. Different works either imposed restrictions on the expressiveness of the policies or were proved secure in the random oracle model. For example, in [20] a bound n was fixed at setup on the number of attributes that could be used while encrypting a message. For constructions that had no bounds on the expressiveness of policies and constant sized public parameters, the random oracle security model was used.

The above restrictions place undesirable burdens on the deployment of ABE schemes. If the designer of the system desires the benefits of avoiding the random oracle heuristic, he has to pick a specific bound for the expressiveness of the system at the setup time; either the size of the attribute universe or the bound on the policies. If the bound is too small, the system might exhaust its functionality and will have to

be completely rebuilt. For example, consider the design of a framework that allows Attribute-Based Encryption in a huge multinational company and suppose that, as this company expands, a large number of new attributes have to be added to the system. If this number exceeds the bound set during the initial deployment of the system, then the company would have to re-deploy the (expanded) system and possibly re-encrypt all its data spending a huge amount of expenses. On the other hand, if the bound chosen is too big, the increased size of the public parameters will impose an unnecessary efficiency burden on all operations.

The first large-universe constructions in the standard model were presented in the recent work of Lewko and Waters [25]. They presented the first large universe KP-ABE construction, secure in the standard model. The system was proved *selectively* secure under static assumptions. The authors of [25] refer to their construction as an “unbounded” scheme, in the sense that the public parameters do not impose additional limitations on the functionality of the systems. Their scheme is indeed large universe, since the size of the attribute universe is exponentially large in the security parameter.

The authors utilized the *dual system* framework on composite order groups to prove security. While this framework is highly useful for the proofs, the actual constructions require use of bilinear groups of large composite order. As a result, these schemes sustain a significant efficiency overhead in comparison to prime order ABE constructions. In a recent result [22] building on [30, 32, 23, 17], one can actually “emulate” the effects of the composite order groups by creating special subspaces of vectors, called *dual vector spaces* introduced by Okamoto and Takashima, and construct a large universe KP-ABE system on prime order groups. This improves the efficiency of the original construction, but there is still a significant performance penalty due to the required size of the vectors.

In a subsequent work by Okamoto and Takashima [33], another unbounded KP-ABE scheme and the first unbounded CP-ABE scheme are presented. Both of them are the first *fully* secure unbounded ABE schemes in the standard model and utilize the *dual vector spaces* framework. Their systems are unbounded, since the size of the public parameters is constant and does not impose any limitations on the size of the attribute universe. However, the traditional notion of “large”-“small” universe is not applicable, since each “attribute” can take one or more different values from an *exponential-sized* space, instead of “present”/“not present”. The total number of these sub-universes should be *polynomial* in the security parameter and the maximum number of times each sub-universe can be used in a policy is a parameter fixed at the setup of the system. As this bound is increased, the overall efficiency of the system is impaired. In our benchmarks, we implement the most efficient “basic version” of their constructions where this parameter is equal to 1, i.e. each sub-universe can be present only once in each policy.

Goals and Contributions We present new constructions and proof techniques for Large Universe ABE in the standard model. Departing from the above recent trends, our constructions are proved *selectively* secure using what is known as partitioning style techniques.

We believe that this is an interesting avenue to explore for two reasons. First, by considering selective model of

security we are able to get more efficient and more practical constructions. While full security is the strongest notion of security, we believe selective is still a meaningful notion and can be a reasonable trade off for performance in some circumstances. In addition, new partitioning proofs can give different and new insights into the security or the style of a construction.

Second, Lewko and Waters [26] recently showed a surprising connection between Dual System Encryption and older selective proofs. Prior fully secure ABE systems [23] required an additional (relative to selective schemes) limit t on the number of times an attribute could be used in a formula. The public parameters and ciphertext size for KP-ABE (key size for CP-ABE) grew proportionally to the bound t . Lewko and Waters showed that through a new “delayed parameter” variant of Dual System Encryption this limit could be done away with. An integral part of their proof was that it leveraged older “program and cancel” style techniques. Given this recent work, a reasonable conclusion is that developing selectively secure proofs might typically become a first step to developing full security. (We note that the large universe construction of [25] was only proved selectively secure.)

We aim to get practical large-universe ABE schemes by adapting and expanding the system from [25] into the prime order setting. In proving security we go back to more traditional “program and cancel” techniques instead of the dual system framework. We present two practical large universe ABE constructions (one CP-ABE and one KP-ABE) in prime order bilinear groups both selectively secure in the standard model under two different q -type assumptions. Our three main objectives in this work were *large universe constructions*, *efficiency*, and *security in the standard model*. Both schemes support a “large universe” attribute space and their public parameters consist of a constant number of group elements. No bounds or other restrictions are imposed on the monotonic Boolean formulas or the attribute sets used by the algorithms of the schemes; thus eliminating the need for design decisions at setup. The efficiency objective refrained us from using composite order groups or dual pairing vector spaces, while to achieve security in the standard model we relied on non static (q -type) assumptions and selective notions. These assumptions are non static in the sense that a polynomial number of terms is given to the adversary and therefore they are intuitively stronger than the static ones. However, the polynomial number of terms gives the ability to the simulator of the proof to embed the additional entropy in the constant number of public parameters. We showcase different techniques for harnessing the power of these assumptions to achieve our large universe constructions. Finally, we demonstrate the efficiency of our constructions by implementing our schemes. We compare performance results to other ABE schemes in prime order groups.

Our Techniques The techniques used to achieve our goals and prove the security of our schemes fall in the category of *partitioning* methodologies. In this setting the simulator of the reduction sets up the public parameters of the systems in such a way that the set of the possible policies (for KP-ABE) or the powerset of the attribute universe (for CP-ABE) is partitioned in two disjoint sets. One for which he can create the secret keys and answer the attackers’ queries, and one for which this is not possible, where the challenge query should belong. Since we are dealing with se-

lective security notions, the simulator knows in advance the required challenge set and therefore the suitable partition. However due to the fact that we are aiming for large universe ABE, which implies constant size public parameters, the simulator has to embed a polynomial amount of “challenge information” in them. This is achieved by utilizing the non static power of our assumptions. Namely, the assumptions’ “size” depends on the size of the declared challenge query. The additional terms available to the simulator allow him to create all the necessary terms for the reduction.

Both our schemes work in a “layered” fashion in order to encrypt information securely and being able to decrypt. In the KP-ABE construction, which is simpler and directly inspired by the composite order construction of [25], two “layers” are employed: the “secret sharing” layer and the “attribute layer”. The first layer is responsible for the sharing of the master secret key during the key generation algorithm and the storing of the blinding factor randomness during the encryption algorithm. The “attribute layer” holds information about the attributes used in both key generation and encryption phases. A “binder term” is utilized to connect the two layers in a secure way. In the CP-ABE construction the situation is slightly more complicated due to the fact that the policies are applied on the ciphertext side. As a result, the “sharing” is applied to the blinding factor randomness and not on the master secret key. Therefore, an additional “binder term” in the public parameters is being used to allow correct decryption using the master secret key. As we will see, the assumptions and the corresponding reductions follow closely this “layer” intuition.

Finally, we mention that both constructions use the “individual randomness” technique from [25] in the “attribute layer” to achieve the large universe functionality. The component for each attribute is masked by a different randomness and as a result no restrictions are imposed on the policies or the attributes, since each component is individually randomized.

1.1 Related Work

Attribute-Based Encryption was introduced by Sahai and Waters [36]. The refinement of the two notions was given in [20] and many CP-ABE and KP-ABE selectively secure constructions followed [6, 15, 19, 34, 35, 43]. Most of them work for monotonic access structures with the exception of the schemes by Ostrovsky, Sahai, and Waters [34], who showed how to realize negation by incorporating specific revocation schemes into the GPSW construction. Fully secure constructions in the standard model were first provided by Okamoto and Takashima [32] and Lewko, Okamoto, Sahai, Takashima, and Waters [23]. The first large universe KP-ABE construction in the standard model was given in [25] (composite order groups) and the first fully secure unbounded constructions were given in [33]. Okamoto and Takashima initiated the dual pairing vector space framework in various works [30, 31, 32], which lead to the first large universe KP-ABE construction in prime order group groups by Lewko [22]. Parameterized (non static) assumptions were introduced in [7] and used in several subsequent works [18, 43]. The problem of an environment with multiple central authorities in ABE was considered in [13, 14, 24], while several authors have presented schemes that do not address the problem of collusion resistance [40, 28, 11, 2, 3, 4].

We note that several techniques in ABE schemes have roots in Identity-Based Encryption [37, 8, 16, 7, 42, 18, 9]. Finally, we mention here the related concept of Predicate Encryption introduced by Katz, Sahai, and Waters [21] and further refined in [39, 38, 31, 23, 32, 10].

1.2 Organization

In Sec. 2 we introduce some notation, background information about access structures and linear secret-sharing schemes, and the complexity assumption for our CP-ABE scheme. Section 3 contains the algorithms and the selective security definition for CP-ABE schemes. Our CP-ABE construction and the security proof are in Sec. 4. Finally, implementations and efficiency results are presented in Sec. 5. The assumption for our KP-ABE construction is in App. A. In App. B and App. C we present the KP-ABE algorithms with the security definition, and our KP-ABE construction with the security proof, respectively.

2. PRELIMINARIES

2.1 Notation

For $n \in \mathbb{N}$, we define $[n] \stackrel{\text{def.}}{=} \{1, 2, \dots, n\}$. Also, for $n_1, n_2, \dots, n_k \in \mathbb{N}$: $[n_1, n_2, \dots, n_k] \stackrel{\text{def.}}{=} [n_1] \times [n_2] \times \dots \times [n_k]$. When \mathcal{S} is a set, we denote by $s \stackrel{\$}{\leftarrow} \mathcal{S}$ the fact that the variable s is picked uniformly at random from \mathcal{S} . We write $s_1, s_2, \dots, s_n \stackrel{\$}{\leftarrow} \mathcal{S}$ as shorthand for $s_1 \stackrel{\$}{\leftarrow} \mathcal{S}, s_2 \stackrel{\$}{\leftarrow} \mathcal{S}, \dots, s_n \stackrel{\$}{\leftarrow} \mathcal{S}$. By $\text{negl}(n)$ we denote a negligible function in n and by PPT probabilistic polynomial-time.

The set of matrices of size $m \times n$ with elements in \mathbb{Z}_p is denoted by $\mathbb{Z}_p^{m \times n}$. Special subsets are the set of row vectors of length n : $\mathbb{Z}_p^{1 \times n}$, and column vectors of length n : $\mathbb{Z}_p^{n \times 1}$. When \vec{v} is a vector (of any type), we will denote by v_i the i -th element and by $\langle \vec{v}, \vec{w} \rangle$ the inner product of vectors \vec{v} and \vec{w} .

2.2 Access Structures and Linear Secret-Sharing Schemes

In this section, we present the formal definitions of access structures and linear secret-sharing schemes introduced in [5], adapted to match our setting.

Definition 2.1 (Access Structures [5]). *Let \mathcal{U} be the attribute universe. An access structure on \mathcal{U} is a collection \mathbb{A} of non-empty sets of attributes, i.e. $\mathbb{A} \subseteq 2^{\mathcal{U}} \setminus \{\emptyset\}$. The sets in \mathbb{A} are called the authorized sets and the sets not in \mathbb{A} are called the unauthorized sets.*

Additionally, an access structure is called monotone if $\forall B, C \in \mathbb{A} : \text{if } B \in \mathbb{A} \text{ and } B \subseteq C, \text{ then } C \in \mathbb{A}$.

In our constructions, we only consider monotone access structures, which means that as a user (CP-ABE setting) acquires more attributes, he will not lose his possible decryption privileges. General access structures in large universe ABE can be realized by splitting the attribute universe in half and treating the attributes of one half as the negated versions of the attributes in the other half [20]. We note that if the access structure is encoded as a monotonic Boolean formula over attributes¹, there is a generic algorithm that

¹A monotonic Boolean formula consists of only AND, OR, and threshold gates, for example $A_1 \wedge (A_2 \vee A_3)$. This means

generates the corresponding access policy in polynomial time [5, 24].

Definition 2.2 (Linear Secret-Sharing Schemes (LSSS) [5]). *Let p be a prime and \mathcal{U} the attribute universe. A secret-sharing scheme Π with domain of secrets \mathbb{Z}_p realizing access structures on \mathcal{U} is linear over \mathbb{Z}_p if*

1. *The shares of a secret $s \in \mathbb{Z}_p$ for each attribute form a vector over \mathbb{Z}_p .*
2. *For each access structure \mathbb{A} on \mathcal{U} , there exists a matrix $M \in \mathbb{Z}_p^{\ell \times n}$, called the share-generating matrix, and a function ρ , that labels the rows of M with attributes from \mathcal{U} , i.e. $\rho: [\ell] \rightarrow \mathcal{U}$, which satisfy the following:*

During the generation of the shares, we consider the column vector $\vec{v} = (s, r_2, \dots, r_n)^\top$, where $r_2, \dots, r_n \xleftarrow{\$} \mathbb{Z}_p$. Then the vector of ℓ shares of the secret s according to Π is equal to $M\vec{v} \in \mathbb{Z}_p^{\ell \times 1}$. The share $(M\vec{v})_j$ where $j \in [\ell]$ “belongs” to attribute $\rho(j)$.

We will be referring to the pair (M, ρ) as the policy of the access structure \mathbb{A} .

According to [5], each secret-sharing scheme (not only the linear ones) should satisfy the *reconstruction requirement* (each authorized set can reconstruct the secret) and the *security requirement* (any unauthorized set cannot reveal any partial information about the secret).

In our setting, let \mathcal{S} denote an authorized set for the access structure \mathbb{A} encoded by the policy (M, ρ) . Then let I be the set of rows whose labels are in \mathcal{S} , i.e. $I = \{i | i \in [\ell] \wedge \rho(i) \in \mathcal{S}\}$. The reconstruction requirement asserts that the vector $(1, 0, \dots, 0)$ is in the span of rows of M indexed by I . This means that there exist constants $\{\omega_i\}_{i \in I}$ in \mathbb{Z}_p such that for any valid shares $\{\lambda_i = (M\vec{v})_i\}_{i \in I}$ of a secret s according to Π , it is true that: $\sum_{i \in I} \omega_i \lambda_i = s$. Additionally, it has been proved in [5] that the constants $\{\omega_i\}_{i \in I}$ can be found in time polynomial in the size of the share-generating matrix M .

On the other hand, for unauthorized sets \mathcal{S}' no such constants $\{\omega_i\}$ exist. Moreover, in this case it is also true that if $I' = \{i | i \in [\ell] \wedge \rho(i) \in \mathcal{S}'\}$, there exists a vector $\vec{w} \in \mathbb{Z}_p^{1 \times n}$, such that its first component w_1 is any non zero element in \mathbb{Z}_p and $\langle \vec{M}_i, \vec{w} \rangle = 0$ for all $i \in I'$, where $\vec{M}_i = (M_{i,1}, M_{i,2}, \dots, M_{i,n})$; the i -th row of M .

2.3 Assumption 1

For our CP-ABE construction we will use a q -type assumption on prime order bilinear groups, denoted by q -1, which is similar to the Decisional Parallel Bilinear Diffie-Hellman Exponent Assumption [43]. It is parameterized by a security parameter $\lambda \in \mathbb{N}$ and an integer q , polynomial in λ . We assume that there exists a group generator algorithm $\mathcal{G}(1^\lambda) \rightarrow (p, \mathbb{G}, \mathbb{G}_T, e)$ that outputs the description of the (symmetric) bilinear group of order $p = \Theta(2^\lambda)$. This assumption can be proved secure in the generic group model, but the proof is omitted due to space constraints. It is defined via the following game between a challenger and an attacker:

Initially the challenger calls the group generation algorithm with input the security parameter, picks a random that as a key (in CP-ABE) or a ciphertext (in KP-ABE) acquires more attributes it will not lose the decryption capabilities.

group element $g \xleftarrow{\$} \mathbb{G}$, and $q+2$ random exponents $a, s, b_1, b_2, \dots, b_q \xleftarrow{\$} \mathbb{Z}_p$. Then he sends to the attacker the group description $(p, \mathbb{G}, \mathbb{G}_T, e)$ and all of the following terms:

$$\begin{aligned} &g, g^s \\ &g^{a^i}, g^{b_j}, g^{sb_j}, g^{a^i b_j}, g^{a^i/b_j^2} \quad \forall (i, j) \in [q, q] \\ &g^{a^i b_j/b_j'^2} \quad \forall (i, j, j') \in [2q, q, q] \text{ with } j \neq j' \\ &g^{a^i/b_j} \quad \forall (i, j) \in [2q, q] \text{ with } i \neq q+1 \\ &g^{sa^i b_j/b_j'}, g^{sa^i b_j/b_j'^2} \quad \forall (i, j, j') \in [q, q, q] \text{ with } j \neq j' \end{aligned}$$

The challenger also flips a random coin $b \xleftarrow{\$} \{0, 1\}$ and if $b = 0$, it gives to the attacker the term $e(g, g)^{sa^{q+1}}$. Otherwise it gives a random term $R \xleftarrow{\$} \mathbb{G}_T$. Finally the attacker outputs a guess $b' \in \{0, 1\}$.

Definition 2.3. *We say that the q -1 assumption holds if all PPT attackers have at most a negligible advantage in λ in the above security game, where the advantage is defined as $\text{Adv} = \Pr[b' = b] - 1/2$.*

Remark: Notice the absence of the term g^{a^{q+1}/b_j} in the fourth line of the assumption. If this term were given to the attacker, then he could break the assumption trivially by pairing it with the corresponding g^{sb_j} term. On the other hand, the term $g^{a^{q+1}b_j/b_j'^2}$ is given, and this poses no problems in the generic group model since $j \neq j'$ and by possible pairing the adversary cannot get rid of the b_j 's.

3. CIPHERTEXT - POLICY ATTRIBUTE - BASED ENCRYPTION

3.1 Algorithms

A Ciphertext-Policy Attribute-Based Encryption scheme consists of the following four PPT algorithms:

- **Setup** $(1^\lambda) \rightarrow (pp, msk)$: The **Setup** algorithm takes the security parameter $\lambda \in \mathbb{N}$ encoded in unary and outputs the public parameters pp and the master secret key msk . We assume that the public parameters contain a description of the attribute universe \mathcal{U} .²

- **KeyGen** $(1^\lambda, pp, msk, \mathcal{S}) \rightarrow sk$: The key generation algorithm takes as inputs the public parameters pp , the master secret key msk and a set of attributes $\mathcal{S} \subseteq \mathcal{U}$. The security parameter is included in the inputs to ensure that it is polynomial time in λ . The algorithm generates a secret key corresponding to \mathcal{S} .

- **Encrypt** $(1^\lambda, pp, m, \mathbb{A}) \rightarrow ct$: The encryption algorithm takes as inputs the public parameters pp , a plaintext message m , and an access structure \mathbb{A} on \mathcal{U} . It outputs the ciphertext ct .

- **Decrypt** $(1^\lambda, pp, sk, ct) \rightarrow m$: The decryption algorithm takes as inputs the public parameters pp , a secret key sk , and a ciphertext ct . It outputs the plaintext m .

Correctness: We require that a CP-ABE scheme is correct, i.e the decryption algorithm correctly decrypts a ciphertext of an access structure \mathbb{A} with a secret key on \mathcal{S} , when \mathcal{S} is an authorized set of \mathbb{A} . Formally:

²In previous CP-ABE constructions the attribute universe \mathcal{U} (or its size) was one of the arguments of the **Setup** algorithm. In our constructions, the attribute universe depends only on the size of the underlying group \mathbb{G} , which depends on the security parameter λ and the group generation algorithm.

Definition 3.1. A CP-ABE scheme is correct when for all messages m , and all attribute sets \mathcal{S} and access structures \mathbb{A} with $\mathcal{S} \in \mathbb{A}$ (i.e. for \mathcal{S} authorized), any pair (pp, msk) output from $\text{Setup}(1^\lambda)$, any secret key sk output from $\text{KeyGen}(1^\lambda, pp, msk, \mathcal{S})$, and any ciphertext ct output by $\text{Encrypt}(1^\lambda, pp, m, \mathbb{A})$, it is true that: $\text{Decrypt}(1^\lambda, pp, sk, ct) = m$.

3.2 CP-ABE Selective Security

In this section we present the definition of selective security for CP-ABE schemes. This is described by a game between a challenger and an attacker and is parameterized by the security parameter $\lambda \in \mathbb{N}$. The phases of the game are the following:

- **Initialization:** In this phase the attacker declares the challenge access structure \mathbb{A}^* , which he will try to attack, and sends it to the challenger.

- **Setup:** Here the challenger calls the $\text{Setup}(1^\lambda)$ algorithm and sends the public parameters pp to the attacker.

- **Query Phase 1:** In this phase the attacker can adaptively ask for secret keys for the sets of attributes $\mathcal{S}_1, \mathcal{S}_2, \dots, \mathcal{S}_{Q_1}$. For each \mathcal{S}_i the challenger calls $\text{KeyGen}(msk, \mathcal{S}_i) \rightarrow sk_i$ and sends sk_i to the attacker. The restriction that has to be satisfied for each query is that none of the queried sets satisfies the challenge access structure, i.e. $\forall i \in [Q_1] : \mathcal{S}_i \notin \mathbb{A}^*$.

- **Challenge:** The attacker declares two equal-length plaintexts m_0 and m_1 and sends them to the challenger. He flips a random coin $b \in \{0, 1\}$ and calls $\text{Encrypt}(m_b, \mathbb{A}^*) \rightarrow ct$. He sends ct to the attacker.

- **Query Phase 2:** This is the same as query phase 1. The attacker asks for the secret key for the sets $\mathcal{S}_{Q_1+1}, \mathcal{S}_{Q_1+2}, \dots, \mathcal{S}_Q$, for which the same restriction holds: $\forall i \in [Q] : \mathcal{S}_i \notin \mathbb{A}^*$.

- **Guess:** The attacker outputs his guess $b' \in \{0, 1\}$ for b .

Definition 3.2. A CP-ABE scheme is selectively secure if all PPT attackers have at most a negligible advantage in λ in the above security game, where the advantage of an attacker is defined as $\text{Adv} = \Pr[b' = b] - 1/2$.

4. OUR LARGE UNIVERSE CP-ABE

In this section we present our large universe CP-ABE construction. The public parameters consist of the six group elements $(g, u, h, w, v, e(g, g)^\alpha)$, which intuitively are utilized in two separate “layers” to achieve secure large universe CP-ABE. In the “attribute layer”, the u, h terms provide a Boneh-Boyen-style [7] hash function $(u^A h)$, while in the “secret sharing layer” the w term holds the secret randomness r during key generation and the shares of the secret randomness s during encryption. The v term is used to “bind” the two layers together. The g and $e(g, g)^\alpha$ terms are used to introduce the master secret key functionality and allow correct decryption.

4.1 Construction

Our scheme consists of the following four algorithms:

- **Setup** $(1^\lambda) \rightarrow (pp, msk)$: The setup algorithm calls the group generator algorithm $\mathcal{G}(1^\lambda)$ and gets the descriptions of the groups and the bilinear mapping $D = (p, \mathbb{G}, \mathbb{G}_T, e)$, where p is the prime order of the groups \mathbb{G} and \mathbb{G}_T . The attribute universe is $\mathcal{U} = \mathbb{Z}_p$.

Then the algorithm picks the random terms $g, u, h, w, v \xleftarrow{\$} \mathbb{G}$ and $\alpha \xleftarrow{\$} \mathbb{Z}_p$. It outputs

$$pp = (D, g, u, h, w, v, e(g, g)^\alpha) \quad msk = (\alpha)$$

- **KeyGen** $(msk, \mathcal{S} = \{A_1, A_2, \dots, A_k\} \subseteq \mathbb{Z}_p) \rightarrow sk$: Initially, the key generation algorithm picks $k+1$ random exponents $r, r_1, r_2, \dots, r_k \xleftarrow{\$} \mathbb{Z}_p$. Then it computes $K_0 = g^\alpha w^r$, $K_1 = g^r$, and for every $\tau \in [k]$

$$K_{\tau,2} = g^{r_\tau} \text{ and } K_{\tau,3} = (u^{A_\tau} h)^{r_\tau} v^{-r}$$

The secret key output is $sk = (\mathcal{S}, K_0, K_1, \{K_{\tau,2}, K_{\tau,3}\}_{\tau \in [k]})$.

- **Encrypt** $(m \in \mathbb{G}_T, (M, \rho)) \rightarrow ct$: The encryption algorithm takes the plaintext message m and the access structure encoded in an LSSS policy, with $M \in \mathbb{Z}_p^{\ell \times n}$ and $\rho : [\ell] \rightarrow \mathbb{Z}_p$. First, it picks $\vec{y} = (s, y_2, \dots, y_n)^\top \xleftarrow{\$} \mathbb{Z}_p^{n \times 1}$. In the terminology of Sec. 2.2, s is the random secret to be shared among the shares. The vector of the shares is $\vec{\lambda} = (\lambda_1, \lambda_2, \dots, \lambda_\ell)^\top = M\vec{y}$.

It then picks ℓ random exponents $t_1, t_2, \dots, t_\ell \xleftarrow{\$} \mathbb{Z}_p$ and calculates $C = m \cdot e(g, g)^{\alpha s}$, $C_0 = g^s$, and for every $\tau \in [\ell]$

$$C_{\tau,1} = w^{\lambda_\tau} v^{t_\tau}, \quad C_{\tau,2} = (u^{\rho(\tau)} h)^{-t_\tau} \text{ and } C_{\tau,3} = g^{t_\tau}$$

The ciphertext output is

$$ct = ((M, \rho), C, C_0, \{C_{\tau,1}, C_{\tau,2}, C_{\tau,3}\}_{\tau \in [\ell]})$$

- **Decrypt** $(sk, ct) \rightarrow m$: Firstly, the decryption algorithm calculates the set of rows in M that provide a share to attributes in \mathcal{S} , i.e. $I = \{i : \rho(i) \in \mathcal{S}\}$. Then it computes the constants $\{\omega_i \in \mathbb{Z}_p\}_{i \in I}$ such that $\sum_{i \in I} \omega_i \vec{M}_i = (1, 0, \dots, 0)$, where \vec{M}_i is the i -th row of the matrix M . These constants exist if the set \mathcal{S} is an authorized set of the policy (c.f. Sec. 2.2).

Then it calculates

$$B = \frac{e(C_0, K_0)}{\prod_{i \in I} (e(C_{i,1}, K_1) e(C_{i,2}, K_{i,2}) e(C_{i,3}, K_{i,3}))^{\omega_i}}$$

where τ is the index of the attribute $\rho(i)$ in \mathcal{S} (it depends on i). The algorithm outputs $m = C/B$.

Correctness: If the attribute set \mathcal{S} of the secret key is authorized, we have that $\sum_{i \in I} \omega_i \lambda_i = s$. Therefore:

$$\begin{aligned} B &= \frac{e(g, g)^{\alpha s} e(g, w)^{rs}}{\prod_{i \in I} e(g, w)^{r \omega_i \lambda_i} e(g, v)^{r t_i \omega_i} e(g, u^{\rho(i)} h)^{-r t_i \omega_i}} \\ &= \frac{1}{\prod_{i \in I} e(g, u^{\rho(i)} h)^{r t_i \omega_i} e(g, v)^{-r t_i \omega_i}} \\ &= \frac{e(g, g)^{\alpha s} e(g, w)^{rs}}{e(g, w)^{r \sum_{i \in I} \omega_i \lambda_i}} = e(g, g)^{\alpha s} \end{aligned}$$

4.2 Proof of Selective Security

We will prove the following theorem regarding the selective security of our CP-ABE scheme:

Theorem 4.1. If the q -1 assumption holds then all PPT adversaries with a challenge matrix of size $\ell \times n$, where $\ell, n \leq q$, have a negligible advantage in selectively breaking our scheme.

Proof. To prove the theorem we will assume that there exists a PPT attacker \mathcal{A} with a challenge matrix that satisfies the restriction, which has a non negligible advantage $\text{Adv}_{\mathcal{A}}$ in selectively breaking our scheme. Using this attacker we will build a PPT simulator \mathcal{B} that attacks the q -1 assumption with a non negligible advantage.

Initialization: \mathcal{B} receives the given terms from the assumption and a challenge policy (M^*, ρ^*) from \mathcal{A} . We have that M^* is an $\ell \times n$ matrix, where $\ell, n \leq q$, and $\rho^* : [\ell] \rightarrow \mathbb{Z}_p$.

Setup: The simulator \mathcal{B} has to provide \mathcal{A} the public parameters of the system. In order to do that it implicitly sets the master secret key of the scheme to be $\alpha = a^{q+1} + \tilde{\alpha}$, where a, q are set in the assumption and $\tilde{\alpha} \xleftarrow{\$} \mathbb{Z}_p$ is a known to \mathcal{B} random exponent. Notice that this way α is correctly distributed and a is information-theoretically hidden from \mathcal{A} . Then \mathcal{B} picks the random exponents $\tilde{v}, \tilde{u}, \tilde{h} \xleftarrow{\$} \mathbb{Z}_p$ and using the assumption gives to \mathcal{A} the following public parameters:

$$\begin{aligned} g &= g \\ u &= g^{\tilde{u}} \cdot \prod_{(j,k) \in [\ell, n]} \left(g^{a^k/b_j^2} \right)^{M_{j,k}^*} \\ h &= g^{\tilde{h}} \cdot \prod_{(j,k) \in [\ell, n]} \left(g^{a^k/b_j^2} \right)^{-\rho^*(j)M_{j,k}^*} \\ w &= g^a \\ v &= g^{\tilde{v}} \cdot \prod_{(j,k) \in [\ell, n]} \left(g^{a^k/b_j} \right)^{M_{j,k}^*} \\ e(g, g)^\alpha &= e(g^a, g^{a^q}) \cdot e(g, g)^{\tilde{\alpha}} \end{aligned}$$

The term w is properly distributed in \mathcal{A} 's view because the term $e(g, g)^\alpha$ hides the exponent a information-theoretically. The terms v, u, h are also properly distributed due to $\tilde{v}, \tilde{u}, \tilde{h}$, respectively. Notice that all terms can be calculated by the simulator using suitable terms from the assumption and the challenge policy given by \mathcal{A} .

As one can see, the “attribute layer”, which consists of the terms u, h , is made up of terms whose exponents have b_i^2 in the denominator, the “binder term” v has b_i , and the “secret sharing layer” w has only one power of a . This scaling of the powers of b_i will allow our simulator to properly simulate all terms.

Query phases 1 and 2: Now the simulator has to produce secret keys for non authorized sets of attributes requested by \mathcal{A} . In both phases the treatment is the same. We describe here the way \mathcal{B} works in order to create a key for an attribute set $\mathcal{S} = \{A_1, A_2, \dots, A_{|\mathcal{S}|}\}$ received by \mathcal{A} .

Since \mathcal{S} is non authorized for (M^*, ρ^*) , there exists a vector $\vec{w} = (w_1, w_2, \dots, w_n)^\top \in \mathbb{Z}_p^n$ such that $w_1 = -1$ and $\langle \vec{M}_i^*, \vec{w} \rangle = 0$ for all $i \in I = \{i | i \in [\ell] \wedge \rho^*(i) \in \mathcal{S}\}$ (c.f. Sec. 2.2). The simulator calculates \vec{w} using linear algebra. Then it picks $\tilde{r} \xleftarrow{\$} \mathbb{Z}_p$ and implicitly sets

$$r = \tilde{r} + w_1 a^q + w_2 a^{q-1} + \dots + w_n a^{q+1-n} = \tilde{r} + \sum_{i \in [n]} w_i a^{q+1-i}$$

This is properly distributed due to \tilde{r} . Then using the suitable terms from the assumption it calculates:

$$\begin{aligned} K_0 &= g^\alpha w^r = g^{a^{q+1}} g^{\tilde{\alpha}} g^{a\tilde{r}} \prod_{i \in [n]} g^{w_i a^{q+2-i}} \\ &= g^{\tilde{\alpha}} (g^a)^{\tilde{r}} \prod_{i=2}^n \left(g^{a^{q+2-i}} \right)^{w_i} \\ K_1 &= g^r = g^{\tilde{r}} \prod_{i \in [n]} \left(g^{a^{q+1-i}} \right)^{w_i} \end{aligned}$$

Additionally, for all $\tau \in [|\mathcal{S}|]$ it has to compute the terms $K_{\tau,2} = g^{r_\tau}$ and $K_{\tau,3} = (u^{A_\tau} h)^{r_\tau} v^{-r}$. The common part v^{-r} for these terms is the following:

$$\begin{aligned} v^{-\tilde{r}} &\left(g^{\tilde{v}} \prod_{(j,k) \in [\ell, n]} g^{a^k M_{j,k}^* / b_j} \right)^{-\sum_{i \in [n]} w_i a^{q+1-i}} \\ &= v^{-\tilde{r}} \prod_{i \in [n]} \left(g^{a^{q+1-i}} \right)^{-\tilde{v} w_i} \cdot \prod_{\substack{(i,j,k) \in \\ [n, \ell, n]}} g^{-w_i M_{j,k}^* a^{q+1+k-i} / b_j} \\ &= v^{-\tilde{r}} \prod_{i \in [n]} \left(g^{a^{q+1-i}} \right)^{-\tilde{v} w_i} \cdot \underbrace{\prod_{\substack{(i,j,k) \in \\ [n, \ell, n], i \neq k}} \left(g^{\frac{a^{q+1+k-i}}{b_j}} \right)^{-w_i M_{j,k}^*}}_{\Phi} \\ &\quad \cdot \prod_{(i,j) \in [n, \ell]} g^{-w_i M_{j,i}^* a^{q+1} / b_j} \\ &= \Phi \cdot \prod_{j \in [\ell]} g^{-\langle \vec{w}, \vec{M}_j^* \rangle a^{q+1} / b_j} = \Phi \cdot \prod_{\substack{j \in [\ell] \\ \rho^*(j) \notin \mathcal{S}}} g^{-\langle \vec{w}, \vec{M}_j^* \rangle a^{q+1} / b_j} \end{aligned}$$

The Φ part can be calculated by the simulator using the assumption, while the second part has to be canceled by the $(u^{A_\tau} h)^{r_\tau}$ part. So for every attribute $A_\tau \in \mathcal{S}$ the simulator sets implicitly

$$\begin{aligned} r_\tau &= \tilde{r}_\tau + r \cdot \sum_{\substack{i' \in [\ell] \\ \rho^*(i') \notin \mathcal{S}}} \frac{b_{i'}}{A_\tau - \rho^*(i')} \\ &= \tilde{r}_\tau + \tilde{r} \cdot \sum_{\substack{i' \in [\ell] \\ \rho^*(i') \notin \mathcal{S}}} \frac{b_{i'}}{A_\tau - \rho^*(i')} + \sum_{\substack{(i,i') \in [n, \ell] \\ \rho^*(i') \notin \mathcal{S}}} \frac{w_i b_{i'} a^{q+1-i}}{A_\tau - \rho^*(i')} \end{aligned}$$

where $\tilde{r}_\tau \xleftarrow{\$} \mathbb{Z}_p$ and therefore r_τ is properly distributed. The use of the b_i 's in the numerators of the fractions is explained by the “layer” intuition presented before. Namely, these b_i will cancel with the b_i^2 denominators in the “attribute layer” and provide a cancellation for the unknown part of v^{-r} .

Also, notice that r_τ is well-defined only for attributes in the specific unauthorized set \mathcal{S} or unrelated attributes (outside the policy), since the sum is over the i' such that $\rho^*(i') \notin \mathcal{S}$. Therefore, for all $A_\tau \in \mathcal{S}$ or $A_\tau \notin \rho^*([\ell])$, the denominators $A_\tau - \rho^*(i')$ are non zero. If the simulator tries to include more attributes of the policy in the key (and possibly make a key for an authorized set), he would have to divide by zero (see Figure 1). Namely, the set of secret keys is partitioned in two sets: the unauthorized, which the simulator can create using the above method, and the authorized, which the simulator cannot create.




Figure 1: The simulator can not create the components for attributes in the gray area.

After some algebraic manipulations, the $(u^{A_\tau} h)^{r_\tau}$ part of $K_{\tau,3}$ is computed as:

$$\begin{aligned}
& (u^{A_\tau} h)^{\tilde{r}_\tau} \cdot (K_{\tau,2}/g^{\tilde{r}_\tau})^{\tilde{u}A_\tau + \tilde{h}} \\
& \cdot \prod_{\substack{(i',j,k) \in [\ell, \ell, n] \\ \rho^*(i') \notin S}} g^{\tilde{r}(A_\tau - \rho^*(j))M_{j,k}^* b_{i'} a^k / (A_\tau - \rho^*(i'))b_j^2} \\
& \cdot \prod_{\substack{(i,i',j,k) \in [n, \ell, \ell, n] \\ \rho^*(i') \notin S}} g^{(A_\tau - \rho^*(j))w_i M_{j,k}^* b_{i'} a^{q+1+k-i} / (A_\tau - \rho^*(i'))b_j^2} \\
& = \Psi \cdot \prod_{\substack{(i,j) \in [n, \ell] \\ \rho^*(j) \notin S}} g^{(A_\tau - \rho^*(j))w_i M_{j,i}^* b_j a^{q+1+i-i} / (A_\tau - \rho^*(j))b_j^2} \\
& = \Psi \cdot \prod_{\substack{j \in [\ell] \\ \rho^*(j) \notin S}} g^{\langle \vec{w}, \vec{M}_j^* \rangle a^{q+1} / b_j}
\end{aligned}$$

where Ψ includes the remaining terms of the product. The Ψ and $K_{\tau,2}$ terms can be calculated using the suitable terms of our assumption. The second part of $(u^{A_\tau} h)^{r_\tau}$ cancels exactly with the problematic part of v^{-r} . Therefore the simulator can calculate $K_{\tau,2}$ and $K_{\tau,3}$ for all $A_\tau \in \mathcal{S}$ and hand over the secret key $sk = (\mathcal{S}, K_0, K_1, \{K_{\tau,2}, K_{\tau,3}\}_{\tau \in [|\mathcal{S}|]})$ to the attacker \mathcal{A} .

Challenge: The attacker will output a pair of messages (m_0, m_1) of the same length. In this phase the simulator flips a random coin $b \xleftarrow{\$} \{0, 1\}$ and constructs

$$C = m_b \cdot T \cdot e(g, g^s)^{\tilde{\alpha}} \quad \text{and} \quad C_0 = g^s$$

where T is the challenge term and g^s the corresponding term of the assumption.

The simulator sets implicitly $\vec{y} = (s, sa + \tilde{y}_2, sa^2 + \tilde{y}_3, \dots, sa^{n-1} + \tilde{y}_n)^\top$, where $\tilde{y}_2, \tilde{y}_3, \dots, \tilde{y}_n \xleftarrow{\$} \mathbb{Z}_p$. We see that the secret s and the vector \vec{y} are properly distributed, since s was information theoretically hidden from \mathcal{A} and the \tilde{y}_i 's are picked uniformly at random. As a result, since $\vec{\lambda} = M^* \vec{y}$ we have that

$$\lambda_\tau = \sum_{i \in [n]} M_{\tau,i}^* sa^{i-1} + \sum_{i=2}^n M_{\tau,i}^* \tilde{y}_i = \sum_{i \in [n]} M_{\tau,i}^* sa^{i-1} + \tilde{\lambda}_\tau$$

for each row $\tau \in [\ell]$.

Notice that the terms $\tilde{\lambda}_\tau = \sum_{i=2}^n M_{\tau,i}^* \tilde{y}_i$ are known to the simulator. For each row the simulator \mathcal{B} sets implicitly $t_\tau = -sb_\tau$. This is properly distributed as well, because the

b_i 's are information theoretically hidden from the attacker. Using the above, \mathcal{B} calculates:

$$\begin{aligned}
C_{\tau,1} &= w^{\lambda_\tau} v^{t_\tau} = w^{\tilde{\lambda}_\tau} \cdot \prod_{i \in [n]} g^{M_{\tau,i}^* sa^i} \cdot (g^{sb_\tau})^{-\tilde{v}} \\
& \cdot \prod_{(j,k) \in [\ell, n]} g^{-M_{j,k}^* a^k sb_\tau / b_j} \\
&= w^{\tilde{\lambda}_\tau} \cdot (g^{sb_\tau})^{-\tilde{v}} \cdot \prod_{i \in [n]} g^{M_{\tau,i}^* sa^i} \cdot \prod_{k \in [n]} g^{-M_{\tau,k}^* a^k sb_\tau / b_\tau} \\
& \cdot \prod_{\substack{(j,k) \in [\ell, n] \\ j \neq \tau}} g^{-M_{j,k}^* a^k sb_\tau / b_j} = \\
&= w^{\tilde{\lambda}_\tau} \cdot (g^{sb_\tau})^{-\tilde{v}} \cdot \prod_{\substack{(j,k) \in [\ell, n] \\ j \neq \tau}} (g^{sa^k b_\tau / b_j})^{-M_{j,k}^*} \\
C_{\tau,2} &= (u^{\rho^*(\tau)} h)^{-t_\tau} = (g^{sb_\tau})^{-(\tilde{u}\rho^*(\tau) + \tilde{h})} \\
& \cdot \left(\prod_{(j,k) \in [\ell, n]} g^{(\rho^*(\tau) - \rho^*(j))M_{j,k}^* a^k / b_j^2} \right)^{-sb_\tau} \\
&= (g^{sb_\tau})^{-(\tilde{u}\rho^*(\tau) + \tilde{h})} \\
& \cdot \prod_{\substack{(j,k) \in [\ell, n] \\ j \neq \tau}} (g^{sa^k b_\tau / b_j^2})^{-(\rho^*(\tau) - \rho^*(j))M_{j,k}^*} \\
C_{\tau,3} &= g^{t_\tau} = (g^{sb_\tau})^{-1}
\end{aligned}$$

Notice that by using $t_\tau = -sb_\tau$ we “raised” the exponents of the “binder” term v so that they cancel with the unknown powers of w^{λ_τ} . Therefore, the simulator hands over the ciphertext $ct = ((M^*, \rho^*), C, C_0, \{C_{\tau,1}, C_{\tau,2}, C_{\tau,3}\}_{\tau \in [\ell]})$ to the attacker \mathcal{A} .

Guess: After the query phase 2, where the simulator creates the secret keys as described above, the attacker outputs a guess b' for the challenge bit. If $b' = b$ the simulator outputs 0, i.e. it claims that the challenge term is $T = e(g, g)^{sa^{q+1}}$. Otherwise, it outputs 1.

If $T = e(g, g)^{sa^{q+1}}$ then \mathcal{A} played the proper security game, because $C = m_b \cdot T \cdot e(g, g^s)^{\tilde{\alpha}} = m_b \cdot e(g, g)^{\alpha s}$. On the other hand, if T is a random term of \mathbb{G}_T then all information about the message m_b is lost in the challenge ciphertext. Therefore the advantage of \mathcal{A} is exactly 0. As a result, if \mathcal{A} breaks the security game with a non negligible advantage, \mathcal{B} has a non negligible advantage in breaking the q -1 assumption. \square

5. IMPLEMENTATION AND EVALUATION

Implementation Details We implemented our schemes in Charm [1]; a framework developed to facilitate the rapid prototyping of cryptographic schemes and protocols. It is based on the Python language which allows the programmer to write code similar to the theoretical implementations. However, the routines that implement the dominant group operations use the PBC library [27] (written natively in C) and the time overhead imposed by the use of Python is

usually less than 1%. Charm also provides routines for applying and using LSSS schemes needed for Attribute-Based systems.

All Charm routines use formally asymmetric groups (although the underlining groups might be symmetric) and therefore we translated our schemes to the asymmetric setting. Namely, we have three groups $\mathbb{G}_1, \mathbb{G}_2$ and \mathbb{G}_T and the pairing e is a function from $\mathbb{G}_1 \times \mathbb{G}_2$ to \mathbb{G}_T . The assumptions and the security proofs can be translated to the asymmetric setting in a generic way. We notice here that we tried to implement our algorithms so that more operations are executed in the \mathbb{G}_1 group than in the \mathbb{G}_2 and that encryption consists mainly of operations in \mathbb{G}_1 , compared to key generation. The reason is that the time taken to execute them in the \mathbb{G}_1 group is considerably smaller than \mathbb{G}_2 in specific asymmetric groups such as the “MNT224” elliptic curve group. We tested the schemes on all ECC groups provided by Charm, i.e. three super-singular symmetric groups and two MNT asymmetric groups [29].

For more information on Charm we refer the reader to [12, 1]. The source code of our implementations can be found in [41]. All our benchmarks were executed on a dual core Intel[®] Xeon[®] CPU W3503@2.40GHz with 2.0GB RAM running Ubuntu R10.04 and Python3.2.3.

We implemented our two ABE schemes (of Sec. 4 and App. C), the prime order KP-ABE construction from [22], and the two basic³ constructions from [33]. Actually, in [22] a large universe prime order HIBE is provided, but the transformation to KP-ABE is straightforward by substituting in the key generation algorithm the additive shares of the secrets with the LSSS shares and the identities with the attributes $\rho(\tau)$. This modified construction is the one we used for comparison to our KP-ABE scheme.

Time Benchmarks In Table 1 we present time benchmarks in different elliptic curve groups for some sample policies (\approx size 4 attributes).

Regarding the comparison between our schemes and prior works, we notice the big gap between the timings of our constructions and prior ones. This is due to the fact that dual vector spaces of high dimension ($\approx 10 - 14$) are utilized, which increase the number of group operations by big factors. We remind the reader that the OT schemes are fully secure, while our schemes and the LW schemes are selectively secure.

Regarding the practicality, in general, of both our schemes we notice that the KeyGen, Encrypt, and Decrypt times of our algorithms are relatively small. They are all under 100ms, with the exception of the super singular 1024-bit curve. Even for this curve the times for each algorithm are under the 700 msec mark. Although one would expect that as the policies and the attributes sets grow bigger these times will increase, the additional overhead will grow only linearly. Thus we believe that the two constructions constitute the most practical implementations of large universe ABE, secure in the standard model.

6. REFERENCES

- [1] Joseph A. Akinyele, Matthew Green, and Avi Rubin. Charm: A framework for rapidly prototyping cryptosystems. Cryptology ePrint Archive, Report 2011/617, 2011. <http://eprint.iacr.org/>.
- [2] Sattam S. Al-Riyami, John Malone-Lee, and Nigel P. Smart. Escrow-free encryption supporting cryptographic workflow. *Int. J. Inf. Sec.*, 5(4):217–229, 2006.
- [3] Walid Bagga, Refik Molva, and Stefano Crosta. Policy-based encryption schemes from bilinear pairings. In *ASIACCS*, page 368, 2006.
- [4] Manuel Barbosa and Pooya Farshim. Secure cryptographic workflow in the standard model. In *INDOCRYPT*, pages 379–393, 2006.
- [5] Amos Beimel. *Secure Schemes for Secret Sharing and Key Distribution*. PhD thesis, Dept. of Computer Science, Technion, 1996.
- [6] John Bethencourt, Amit Sahai, and Brent Waters. Ciphertext-policy attribute-based encryption. In *IEEE Symposium on Security and Privacy*, pages 321–334, 2007.
- [7] Dan Boneh and Xavier Boyen. Efficient selective-id secure identity-based encryption without random oracles. In *EUROCRYPT*, pages 223–238, 2004.
- [8] Dan Boneh and Matthew K. Franklin. Identity-based encryption from the Weil pairing. In *CRYPTO*, pages 213–229, 2001.
- [9] Dan Boneh, Craig Gentry, and Michael Hamburg. Space-efficient identity based encryption without pairings. In *FOCS*, pages 647–657, 2007.
- [10] Dan Boneh, Amit Sahai, and Brent Waters. Functional encryption: Definitions and challenges. In *TCC*, pages 253–273, 2011.
- [11] Robert W. Bradshaw, Jason E. Holt, and Kent E. Seamons. Concealing complex policies with hidden credentials. In *ACM Conference on Computer and Communications Security*, pages 146–157, 2004.
- [12] Charm. <http://www.charm-crypto.com>.
- [13] Melissa Chase. Multi-authority attribute based encryption. In *TCC*, pages 515–534, 2007.
- [14] Melissa Chase and Sherman S. M. Chow. Improving privacy and security in multi-authority attribute-based encryption. In *ACM Conference on Computer and Communications Security*, pages 121–130, 2009.
- [15] Ling Cheung and Calvin C. Newport. Provably secure ciphertext policy ABE. In *ACM Conference on Computer and Communications Security*, pages 456–465, 2007.
- [16] Clifford Cocks. An identity based encryption scheme based on quadratic residues. In *IMA Int. Conf.*, pages 360–363, 2001.
- [17] David Mandell Freeman. Converting pairing-based cryptosystems from composite-order groups to prime-order groups. In *EUROCRYPT*, pages 44–61, 2010.
- [18] Craig Gentry. Practical identity-based encryption without random oracles. In *EUROCRYPT*, pages 445–464, 2006.
- [19] Vipul Goyal, Abhishek Jain, Omkant Pandey, and Amit Sahai. Bounded ciphertext policy attribute based encryption. In *ICALP (2)*, pages 579–591, 2008.
- [20] Vipul Goyal, Omkant Pandey, Amit Sahai, and Brent Waters. Attribute-based encryption for fine-grained

[1] Joseph A. Akinyele, Matthew Green, and Avi Rubin. Charm: A framework for rapidly prototyping

³Where each sub-universe can appear at most once in the policy.

Curve	Type	Scheme	Setup	KeyGen	Encrypt	Decrypt
“SS512”	KP-ABE	[App.C]	19.1	49.1	30.7	14.7
		LW [22]	447.2	642.3	483.4	44.7
		OT [33]	673.7	924.4	933.5	65.6
	CP-ABE	[Sec.4] OT [33]	25.0 678.0	32.9 922.9	52.0 938.5	16.6 66.0
“SS1024”	KP-ABE	[App.C]	71.5	626.3	396.8	325.3
		LW [22]	5553.3	9283.8	6978.3	1098.8
		OT [33]	7904.3	13389.3	13582.0	1735.7
	CP-ABE	[Sec.4] OT [33]	110.8 7898.9	431.0 13393.2	669.3 13598.7	374.4 1740.4
“MNT159”	KP-ABE	[App.C]	21.1	48.1	44.3	36.4
		LW [22]	692.2	1666.3	168.9	125.2
		OT [33]	930.7	2435.1	320.6	178.4
	CP-ABE	[Sec.4] OT [33]	23.5 929.9	43.8 2396.2	53.5 326.7	41.5 183.5
“MNT201”	KP-ABE	[App.C]	28.4	59.2	60.2	49.7
		LW [22]	929.8	2301.1	237.8	173.6
		OT [33]	1237.1	3338.3	453.3	251.5
	CP-ABE	[Sec.4] OT [33]	31.3 1235.1	58.7 3328.7	71.9 463.3	57.4 251.8
“MNT224”	KP-ABE	[App.C]	34.2	73.4	74.2	60.9
		LW [22]	1150.9	2896.0	302.1	215.6
		OT [33]	1514.9	4156.3	572.4	309.8
	CP-ABE	[Sec.4] OT [33]	37.9 1511.7	73.2 4140.0	88.2 584.5	74.4 310.7

Table 1: Typical running times in milliseconds of each scheme. KeyGen and Encrypt are called with attribute sets and policies of size 4, while Decrypt with common attribute sets of size 2. “MNT” are the Miyaji, Nakabayashi, Takano curves (asymmetric pairing groups), while “SS” are super singular curves (symmetric pairing groups). The number after the type of the curve denotes the size of the base field in bits.

- access control of encrypted data. In *ACM Conference on Computer and Communications Security*, pages 89–98, 2006.
- [21] Jonathan Katz, Amit Sahai, and Brent Waters. Predicate encryption supporting disjunctions, polynomial equations, and inner products. In *EUROCRYPT*, pages 146–162, 2008.
- [22] Allison B. Lewko. Tools for simulating features of composite order bilinear groups in the prime order setting. In *EUROCRYPT*, pages 318–335, 2012.
- [23] Allison B. Lewko, Tatsuaki Okamoto, Amit Sahai, Katsuyuki Takashima, and Brent Waters. Fully secure functional encryption: Attribute-based encryption and (hierarchical) inner product encryption. In *EUROCRYPT*, pages 62–91, 2010.
- [24] Allison B. Lewko and Brent Waters. Decentralizing attribute-based encryption. In *EUROCRYPT*, pages 568–588, 2011.
- [25] Allison B. Lewko and Brent Waters. Unbounded HIBE and attribute-based encryption. In *EUROCRYPT*, pages 547–567, 2011.
- [26] Allison B. Lewko and Brent Waters. New proof methods for attribute-based encryption: Achieving full security through selective techniques. In *CRYPTO*, pages 180–198, 2012.
- [27] Ben Lynn. The Stanford pairing based crypto library. <http://crypto.stanford.edu/abc>.
- [28] Gerome Miklau and Dan Suciu. Controlling access to published data using cryptography. In *VLDB*, pages 898–909, 2003.
- [29] Atsuko Miyaji, Masaki Nakabayashi, and Shunzo Takano. Characterization of elliptic curve traces under fr-reduction. In *ICISC*, pages 90–108, 2000.
- [30] Tatsuaki Okamoto and Katsuyuki Takashima. Homomorphic encryption and signatures from vector decomposition. In *Pairing*, pages 57–74, 2008.
- [31] Tatsuaki Okamoto and Katsuyuki Takashima. Hierarchical predicate encryption for inner-products. In *ASIACRYPT*, pages 214–231, 2009.
- [32] Tatsuaki Okamoto and Katsuyuki Takashima. Fully secure functional encryption with general relations from the decisional linear assumption. In *CRYPTO*, pages 191–208, 2010.
- [33] Tatsuaki Okamoto and Katsuyuki Takashima. Fully secure unbounded inner-product and attribute-based encryption. In *ASIACRYPT*, pages 349–366, 2012.
- [34] Rafail Ostrovsky, Amit Sahai, and Brent Waters. Attribute-based encryption with non-monotonic access structures. In *ACM Conference on Computer and Communications Security*, pages 195–203, 2007.
- [35] Matthew Pirretti, Patrick Traynor, Patrick McDaniel, and Brent Waters. Secure attribute-based systems. In *ACM Conference on Computer and Communications Security*, pages 99–112, 2006.

- [36] Amit Sahai and Brent Waters. Fuzzy identity-based encryption. In *EUROCRYPT*, pages 457–473, 2005.
- [37] Adi Shamir. Identity-based cryptosystems and signature schemes. In *CRYPTO*, pages 47–53, 1984.
- [38] Emily Shen, Elaine Shi, and Brent Waters. Predicate privacy in encryption systems. In *TCC*, pages 457–473, 2009.
- [39] Elaine Shi and Brent Waters. Delegating capabilities in predicate encryption systems. In *ICALP (2)*, pages 560–578, 2008.
- [40] Nigel P. Smart. Access control using pairing based cryptography. In *CT-RSA*, pages 111–121, 2003.
- [41] Source code of our constructions.
www.cs.utexas.edu/~jrous/.
- [42] Brent Waters. Efficient identity-based encryption without random oracles. In *EUROCRYPT*, pages 114–127, 2005.
- [43] Brent Waters. Ciphertext-policy attribute-based encryption: An expressive, efficient, and provably secure realization. In *Public Key Cryptography*, pages 53–70, 2011.

APPENDIX

A. ASSUMPTION 2

For our KP-ABE construction we will use a q -type assumption on prime order bilinear groups, denoted by q -2, which is similar to the Decisional Bilinear Diffie-Hellman Assumption augmented with q parameters b_i . It is parameterized by a security parameter $\lambda \in \mathbb{N}$ and an integer q , polynomial in λ . We assume that there exists a group generator algorithm $\mathcal{G}(1^\lambda) \rightarrow (p, \mathbb{G}, \mathbb{G}_T, e)$ that outputs the description of the (symmetric) bilinear group of order $p = \Theta(2^\lambda)$. This assumption can be proved secure in the generic group model, but the proof is omitted due to space constraints. It is defined via the following game between a challenger and an attacker:

Initially the challenger calls the group generation algorithm with input the security parameter, picks a random group element $g \xleftarrow{\$} \mathbb{G}$, and $q + 3$ random exponents $x, y, z, b_1, b_2, \dots, b_q \xleftarrow{\$} \mathbb{Z}_p$. Then he sends to the attacker the group description $(p, \mathbb{G}, \mathbb{G}_T, e)$ and all of the following terms:

$$\begin{aligned} &g, g^x, g^y, g^z, g^{(xz)^2} \\ &g^{b_i}, g^{xz b_i}, g^{xz/b_i}, g^{x^2 z b_i}, g^{y/b_i^2}, g^{y^2/b_i^2} \quad \forall i \in [q] \\ &g^{xz b_i/b_j}, g^{y b_i/b_j^2}, g^{xy z b_i/b_j}, g^{(xz)^2 b_i/b_j} \quad \forall i, j \in [q], i \neq j \end{aligned}$$

The challenger also flips a random coin $b \xleftarrow{\$} \{0, 1\}$ and if $b = 0$ it gives to the attacker the term $e(g, g)^{xyz}$. Otherwise it gives a random term $R \xleftarrow{\$} \mathbb{G}_T$. Finally the attacker outputs a guess $b' \in \{0, 1\}$.

Definition A.1. We say that the q -2 assumption holds if all PPT attackers have at most a negligible advantage in λ in the above security game, where the advantage is defined as $\text{Adv} = \Pr[b' = b] - 1/2$.

B. KEY-POLICY ATTRIBUTE-BASED ENCRYPTION

B.1 Algorithms

A Key-Policy Attribute-Based Encryption scheme consists of the following four PPT algorithms:

- **Setup**(1^λ) $\rightarrow (pp, msk)$: The **Setup** algorithm takes the security parameter $\lambda \in \mathbb{N}$ encoded in unary and outputs the public parameters pp and the master secret key msk . We assume that the public parameters contain a description of the attribute universe \mathcal{U} .
- **KeyGen**($1^\lambda, pp, msk, \mathbb{A}$) $\rightarrow sk$: The key generation algorithm takes as inputs the public parameters pp , the master secret key msk and an access structure \mathbb{A} on \mathcal{U} . The algorithm generates a secret key corresponding to \mathbb{A} .
- **Encrypt**($1^\lambda, pp, m, S$) $\rightarrow ct$: The encryption algorithm takes as inputs the public parameters pp , a plaintext message m , and a set of attributes $S \subseteq \mathcal{U}$. It outputs the ciphertext ct .
- **Decrypt**($1^\lambda, pp, sk, ct$) $\rightarrow m$: The decryption algorithm takes as inputs the public parameters pp , a secret key sk , and a ciphertext ct . It outputs the plaintext m .

Correctness: We require that a KP-ABE scheme is correct, i.e the decryption algorithm correctly decrypts a ciphertext on S with a secret key of an access structure \mathbb{A} when S is an authorized set of \mathbb{A} . Formally:

Definition B.1. A KP-ABE scheme is correct when for all messages m , and all attribute sets S and access structures \mathbb{A} with $S \in \mathbb{A}$ (i.e. for S authorized), any pair (pp, msk) output from **Setup**(1^λ), any secret key sk output from **KeyGen**($1^\lambda, pp, msk, \mathbb{A}$), and any ciphertext ct output by **Encrypt**($1^\lambda, pp, m, S$), it is true that: **Decrypt**($1^\lambda, pp, sk, ct$) = m .

B.2 KP-ABE Selective Security

The selective security game for KP-ABE is described by a game between a challenger and an attacker and is parameterized by the security parameter $\lambda \in \mathbb{N}$. The phases of the game are the following:

- **Initialization:** In this phase the attacker declares the challenge attribute set S^* , which he will try to attack, and sends it to the challenger.
- **Setup:** Here the challenger calls the **Setup**(1^λ) algorithm and sends the public parameters pp to the attacker.
- **Query Phase 1:** In this phase the attacker can adaptively ask for secret keys for the access structures $\mathbb{A}_1, \mathbb{A}_2, \dots, \mathbb{A}_{Q_1}$. For each \mathbb{A}_i the challenger calls **KeyGen**(msk, \mathbb{A}_i) $\rightarrow sk_i$ and sends sk_i to the attacker. The restriction that has to be satisfied for each query is that none of the queried policies is satisfied by the challenge attribute set, i.e. $\forall i \in [Q_1] : S^* \notin \mathbb{A}_i$.
- **Challenge:** The attacker declares two equal-length plaintexts m_0 and m_1 and sends them to the challenger. He flips a random coin $b \in \{0, 1\}$ and calls **Encrypt**(m_b, S^*) $\rightarrow ct$. He sends ct to the attacker.
- **Query Phase 2:** This is the same as query phase 1. The attacker asks for the secret key for the access structures $\mathbb{A}_{Q_1+1}, \mathbb{A}_{Q_1+2}, \dots, \mathbb{A}_Q$, for which the same restriction holds: $\forall i \in [Q] : S^* \notin \mathbb{A}_i$.
- **Guess:** The attacker outputs his guess $b' \in \{0, 1\}$ for b .

Definition B.2. A KP-ABE scheme is selectively secure if all PPT attackers have at most a negligible advantage in λ in the above security game, where the advantage of an attacker is defined as $\text{Adv} = \Pr[b' = b] - 1/2$.

C. OUR LARGE UNIVERSE KP-ABE

In this section we present our large universe KP-ABE scheme. We mention here that it can be converted to an HIBE scheme using non repeating identities, “AND” policies and delegation capabilities (c.f. [25]). The intuition behind the functionality of this construction is simpler than the CP-ABE. In this setting the public parameters consist of the five terms $(g, u, h, w, e(g, g)^\alpha)$. There is one term less due to the fact that now the master secret key α is the secret to be shared during all the key generation calls. As a result the “secret sharing layer” uses the g term only and the w term is used to “bind” this layer to the u, h “attribute layer”.

C.1 Construction

Our scheme consists of the following four algorithms.

- **Setup** $(1^\lambda) \rightarrow (pp, msk)$: The setup algorithm calls the group generator algorithm $\mathcal{G}(1^\lambda)$ and gets the descriptions of the groups and the bilinear mapping $D = (p, \mathbb{G}, \mathbb{G}_T, e)$, where p is the prime order of the groups \mathbb{G} and \mathbb{G}_T . The attribute universe is $\mathcal{U} = \mathbb{Z}_p$.

Then the algorithm picks the random terms $g, u, h, w \xleftarrow{\$} \mathbb{G}$ and $\alpha \xleftarrow{\$} \mathbb{Z}_p$. It outputs

$$pp = (D, g, u, h, w, e(g, g)^\alpha) \quad msk = (\alpha)$$

- **KeyGen** $(msk, (M, \rho)) \rightarrow sk$: Initially, the algorithm picks $\vec{y} = (\alpha, y_2, \dots, y_n)^\top$ where $y_2, \dots, y_n \xleftarrow{\$} \mathbb{Z}_p$. In the terminology of Sec. 2.2, the master secret key α is the secret to be shared among the shares. The vector of the shares is

$$\vec{\lambda} = (\lambda_1, \lambda_2, \dots, \lambda_\ell)^\top = M\vec{y}$$

It then picks ℓ random exponents $t_1, t_2, \dots, t_\ell \xleftarrow{\$} \mathbb{Z}_p$ and for every $\tau \in [\ell]$ it computes

$$K_{\tau,0} = g^{\lambda_\tau} w^{t_\tau} \quad K_{\tau,1} = \left(u^{\rho(\tau)} h\right)^{-t_\tau} \quad K_{\tau,2} = g^{t_\tau}$$

The secret key is $sk = ((M, \rho), \{K_{\tau,0}, K_{\tau,1}, K_{\tau,2}\}_{\tau \in [\ell]})$.

- **Encrypt** $(m, \mathcal{S} = \{A_1, A_2, \dots, A_k\} \subseteq \mathbb{Z}_p) \rightarrow ct$: Initially, the algorithm picks $k+1$ random exponents $s, r_1, r_2, \dots, r_k \xleftarrow{\$} \mathbb{Z}_p$. It computes $C = m \cdot e(g, g)^{\alpha s}$, $C_0 = g^s$, and for every $\tau \in [k]$ it computes

$$C_{\tau,1} = g^{r_\tau} \quad C_{\tau,2} = (u^{A_\tau} h)^{r_\tau} w^{-s}$$

The ciphertext is $ct = (\mathcal{S}, C, C_0, \{C_{\tau,1}, C_{\tau,2}\}_{\tau \in [k]})$.

- **Decrypt** $(sk, ct) \rightarrow m$: The algorithm finds the set of rows in M that provide a share to attributes in \mathcal{S} , i.e. $I = \{i : \rho(i) \in \mathcal{S}\}$. Then it calculates constants $\{\omega_i \in \mathbb{Z}_p\}_{i \in I}$ such that $\sum_{i \in I} \omega_i \vec{M}_i = (1, 0, \dots, 0)$, where M_i is the i -th row of the matrix M . These constants exist if the set \mathcal{S} is an authorized set of the policy (c.f. Sec. 2.2).

Then it calculates

$$B = \prod_{i \in I} (e(C_0, K_{i,0}) e(C_{\tau,1}, K_{i,1}) e(C_{\tau,2}, K_{i,2}))^{\omega_i}$$

where τ is the index of the attribute $\rho(i)$ in \mathcal{S} (it depends on i). The algorithm outputs $m = C/B$.

Correctness: If the attribute set \mathcal{S} of the ciphertext is authorized, we have that $\sum_{i \in I} \omega_i \lambda_i = \alpha$. Therefore:

$$\begin{aligned} B &= \prod_{i \in I} e(g, g)^{s \omega_i \lambda_i} e(g, w)^{s t_i \omega_i} \\ &\quad \cdot e(g, u^{\rho(i)} h)^{-r_\tau t_i \omega_i} e(g, u^{\rho(i)} h)^{r_\tau t_i \omega_i} e(g, w)^{-s t_i \omega_i} \\ &= e(g, g)^{s \sum_{i \in I} \omega_i \lambda_i} = e(g, g)^{\alpha s} \end{aligned}$$

C.2 Proof of Selective Security

We will prove the following theorem regarding the selective security of our KP-ABE scheme:

Theorem C.1. If the q -2 assumption holds, then all PPT adversaries with a challenge attribute set of size k , where $k \leq q$, have a negligible advantage in selectively breaking our scheme.

Proof. To prove the theorem we will assume that there exists a PPT attacker \mathcal{A} with a challenge attribute set that satisfies the restriction, which has a non negligible advantage $\text{Adv}_{\mathcal{A}}$ in selectively breaking our scheme. Using this attacker we will build a PPT simulator \mathcal{B} that attacks the q -2 assumption with a non negligible advantage.

Initialization: Initially, \mathcal{B} receives the given terms from the assumption and an attribute set $\mathcal{S}^* = \{A_1^*, A_2^*, \dots, A_k^*\} \subseteq \mathcal{U}$.

Setup: Now, the simulator \mathcal{B} has to provide \mathcal{A} the public parameters of the system. In order to do that it implicitly sets the master secret key of the scheme to be $\alpha = xy$, where x, y are set in the assumption. Notice that this way α is properly distributed. Then \mathcal{B} picks the random exponents $\tilde{u}, \tilde{h} \xleftarrow{\$} \mathbb{Z}_p$ and gives to \mathcal{A} the following terms:

$$\begin{aligned} g &= g \\ u &= g^{\tilde{u}} \cdot \prod_{i \in [k]} g^{y/b_i^2} \\ h &= g^{\tilde{h}} \cdot \prod_{i \in [k]} g^{x z/b_i} \cdot \prod_{i \in [k]} \left(g^{y/b_i^2}\right)^{-A_i^*} \\ w &= g^x \\ e(g, g)^\alpha &= e(g^x, g^y) \end{aligned}$$

Since x is information-theoretically hidden from \mathcal{A} , because it is multiplied by y in α , the term w is properly uniformly random in \mathbb{G} . The terms u, h are properly distributed due to \tilde{u}, \tilde{h} respectively. Notice that all terms can be calculated by the simulator using suitable terms from the assumption and the challenge set \mathcal{S}^* given by \mathcal{A} .

In the KP-ABE proof we see that the “binder term” of the CP-ABE reduction has been contained in the “attribute layer”; namely the g^{xz/b_i} of the h term. Since the master secret key α is shared in all key generation queries in KP-ABE, we don’t need any more the extra functionality provided by the powers of a .

Query phases 1 and 2: The simulator has to produce secret keys for policies requested by \mathcal{A} , for which the set \mathcal{S}^* is not authorized. In both phases the treatment is the same. We describe here the way \mathcal{B} works in order to create a key for a policy (M, ρ) with $M \in \mathbb{Z}_p^{\ell \times n}$ and $\rho : [\ell] \rightarrow \mathbb{Z}_p$.

Since \mathcal{S}^* is not authorized for (M, ρ) , there exists a vector $\vec{w} = (w_1, w_2, \dots, w_n)^\top \in \mathbb{Z}_p^n$ such that $w_1 = 1$ and $\langle \vec{M}_\tau, \vec{w} \rangle = 0$ for all $\tau \in [\ell]$ such that $\rho(\tau) \in \mathcal{S}^*$ (c.f. Sec. 2.2). The simulator calculates \vec{w} using linear algebra. The

vector \vec{y} that will be shared is implicitly

$$\vec{y} = xy\vec{w} + (0, \tilde{y}_2, \tilde{y}_3, \dots, \tilde{y}_n)^\top$$

where $\tilde{y}_2, \tilde{y}_3, \dots, \tilde{y}_n \xleftarrow{\$} \mathbb{Z}_p$. This vector is properly distributed because its first component is $xy = \alpha$ and the remaining components are uniformly random in \mathbb{Z}_p . Therefore for each row $\tau \in [\ell]$ the share is

$$\begin{aligned} \lambda_\tau &= \langle \vec{M}_\tau, \vec{y} \rangle = xy \langle \vec{M}_\tau, \vec{w} \rangle + \langle \vec{M}_\tau, (0, \tilde{y}_2, \tilde{y}_3, \dots, \tilde{y}_n)^\top \rangle \\ &= xy \langle \vec{M}_\tau, \vec{w} \rangle + \tilde{\lambda}_\tau \end{aligned}$$

As we mentioned above for each row τ for which $\rho(\tau) \in \mathcal{S}^*$ it is true that $\langle \vec{M}_\tau, \vec{w} \rangle = 0$. Therefore in this case $\lambda_\tau = \tilde{\lambda}_\tau = \langle \vec{M}_\tau, (0, \tilde{y}_2, \tilde{y}_3, \dots, \tilde{y}_n)^\top \rangle$; hence its value is known to the simulator. In that case it picks $t_\tau \xleftarrow{\$} \mathbb{Z}_p$ and outputs the terms $K_{\tau,0}, K_{\tau,1}, K_{\tau,2}$ as in the **KeyGen** algorithm.

On the other hand, for each row τ for which $\rho(\tau) \notin \mathcal{S}^*$ it picks $\tilde{t}_\tau \xleftarrow{\$} \mathbb{Z}_p$ and sets implicitly

$$t_\tau = -y \langle \vec{M}_\tau, \vec{w} \rangle + \sum_{i \in [k]} \frac{xzb_i \langle \vec{M}_\tau, \vec{w} \rangle}{\rho(\tau) - A_i^*} + \tilde{t}_\tau$$

Since $\rho(\tau) \notin \mathcal{S}^*$ the above fractions are defined and t_τ is properly distributed due to \tilde{t}_τ . The intuition behind this choice is that the y exponent “raises” the power of w to the secret $\alpha = xy$. However, this also results to xyz/b_i exponents from h . Thus, the cancellation is provided by the xzb_i exponents on the y/b_i^2 part. Now the simulator can compute the following terms using the assumption:

$$\begin{aligned} K_{\tau,0} &= g^{\lambda_\tau} w^{t_\tau} \\ &= g^{xy \langle \vec{M}_\tau, \vec{w} \rangle + \tilde{\lambda}_\tau} \cdot g^{-xy \langle \vec{M}_\tau, \vec{w} \rangle + \sum_{i \in [k]} \frac{xzb_i \langle \vec{M}_\tau, \vec{w} \rangle}{\rho(\tau) - A_i^*}} \cdot w^{\tilde{t}_\tau} \\ &= g^{\tilde{\lambda}_\tau} \cdot \prod_{i \in [n]} \left(g^{xzb_i} \right)^{\langle \vec{M}_\tau, \vec{w} \rangle / (\rho(\tau) - A_i^*)} \cdot w^{\tilde{t}_\tau} \\ K_{\tau,1} &= (u^{\rho(\tau)} h)^{-t_\tau} = \\ &= g^{y \langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) \tilde{u} + \tilde{h})} \prod_{i \in [k]} g^{-xzb_i (\rho(\tau) \tilde{u} + \tilde{h}) \langle \vec{M}_\tau, \vec{w} \rangle / (\rho(\tau) - A_i^*)} \\ &\quad \cdot \prod_{i \in [k]} g^{xyz \langle \vec{M}_\tau, \vec{w} \rangle / b_i} \prod_{(i,j) \in [k,k]} g^{-(xz)^2 b_j \langle \vec{M}_\tau, \vec{w} \rangle / b_i (\rho(\tau) - A_j^*)} \\ &\quad \cdot \prod_{i \in [k]} g^{y^2 \langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) - A_i^*) / b_i^2} \cdot (u^{\rho(\tau)} h)^{-\tilde{t}_\tau} \\ &\quad \cdot \prod_{(i,j) \in [k,k]} g^{-xyz \langle \vec{M}_\tau, \vec{w} \rangle b_j (\rho(\tau) - A_i^*) / b_i^2 (\rho(\tau) - A_j^*)} \\ &= (g^y)^{\langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) \tilde{u} + \tilde{h})} \prod_{i \in [k]} \left(g^{xzb_i} \right)^{-\langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) \tilde{u} + \tilde{h}) \langle \vec{M}_\tau, \vec{w} \rangle / (\rho(\tau) - A_i^*)} \\ &\quad \cdot \prod_{(i,j) \in [k,k]} \left(g^{(xz)^2 b_j / b_i} \right)^{-\langle \vec{M}_\tau, \vec{w} \rangle / (\rho(\tau) - A_j^*)} \\ &\quad \cdot \prod_{i \in [k]} \left(g^{y^2 / b_i^2} \right)^{\langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) - A_i^*)} \cdot (u^{\rho(\tau)} h)^{-\tilde{t}_\tau} \\ &\quad \cdot \prod_{\substack{(i,j) \in [k,k] \\ i \neq j}} \left(g^{xyz b_j / b_i^2} \right)^{-\langle \vec{M}_\tau, \vec{w} \rangle (\rho(\tau) - A_i^*) / (\rho(\tau) - A_j^*)} \end{aligned}$$

$$K_{\tau,2} = g^{t_\tau}$$

$$= (g^y)^{-\langle \vec{M}_\tau, \vec{w} \rangle} \cdot \prod_{i \in [k]} \left(g^{xzb_i} \right)^{\langle \vec{M}_\tau, \vec{w} \rangle / (\rho(\tau) - A_i^*)} \cdot g^{\tilde{t}_\tau}$$

Therefore \mathcal{B} can reply to \mathcal{A} 's query with the entire secret key $sk = ((M, \rho), \{K_{\tau,0}, K_{\tau,1}, K_{\tau,2}\}_{\tau \in [\ell]})$.

Challenge: The attacker will output a pair of messages (m_0, m_1) of the same length. In this phase the simulator flips a random coin $b \xleftarrow{\$} \{0, 1\}$ and sets implicitly $s = z$ from the q -2 assumption. Also, it sets $r_\tau = b_\tau$ for every level $\tau \in [k]$. These parameters are properly distributed since z, b_1, \dots, b_q are information-theoretically hidden from the attacker's view. Now the simulator can compute the following terms using the assumption:

$$\begin{aligned} C &= m_b \cdot T \quad C_0 = g^s = g^z \\ C_{\tau,1} &= g^{r_\tau} = g^{b_\tau} \\ C_{\tau,2} &= (u^{A_\tau^*} h)^{r_\tau} \cdot w^{-s} \\ &= g^{b_\tau (\tilde{u} A_\tau^* + \tilde{h})} \cdot \prod_{i \in [k]} g^{xzb_\tau / b_i} \prod_{i \in [k]} g^{yb_\tau (A_k^* - A_i^*) / b_i^2} \cdot g^{-xz} \\ &= \left(g^{b_\tau} \right)^{\tilde{u} A_\tau^* + \tilde{h}} \cdot \prod_{\substack{i \in [k] \\ i \neq \tau}} g^{xzb_\tau / b_i} \prod_{\substack{i \in [k] \\ i \neq \tau}} \left(g^{yb_\tau / b_i^2} \right)^{A_\tau^* - A_i^*} \end{aligned}$$

As one can see, the choice of $r_\tau = b_\tau$ “raises” one of the xz/b_i components to xz and achieves the cancellation with w^{-s} . The simulator hands over the ciphertext $ct = (\mathcal{S}^*, C, C_0, \{C_{\tau,1}, C_{\tau,2}\}_{\tau \in [k]})$ to the attacker \mathcal{A} .

Guess: After the query phase 2, where the simulator creates the secret keys as described above, the attacker outputs a guess b' for the challenge bit. If $b' = b$ the simulator outputs 0, i.e. it claims that the challenge term is $T = e(g, g)^{xyz}$. Otherwise, it outputs 1.

If $T = e(g, g)^{xyz}$ then \mathcal{A} played the proper security game, because $C = m_b \cdot T = m_b \cdot e(g, g)^{\alpha s}$. On the other hand, if T is a random term of \mathbb{G}_T then all information about the message m_b is lost in the challenge ciphertext. Therefore the advantage of \mathcal{A} is exactly 0. As a result, if \mathcal{A} breaks the security game with a non negligible advantage, \mathcal{B} has a non negligible advantage in breaking the q -2 assumption. \square